

---

# Deep Learning Model Optimization

## Quantization and Pruning for Enhanced Efficiency and Performance

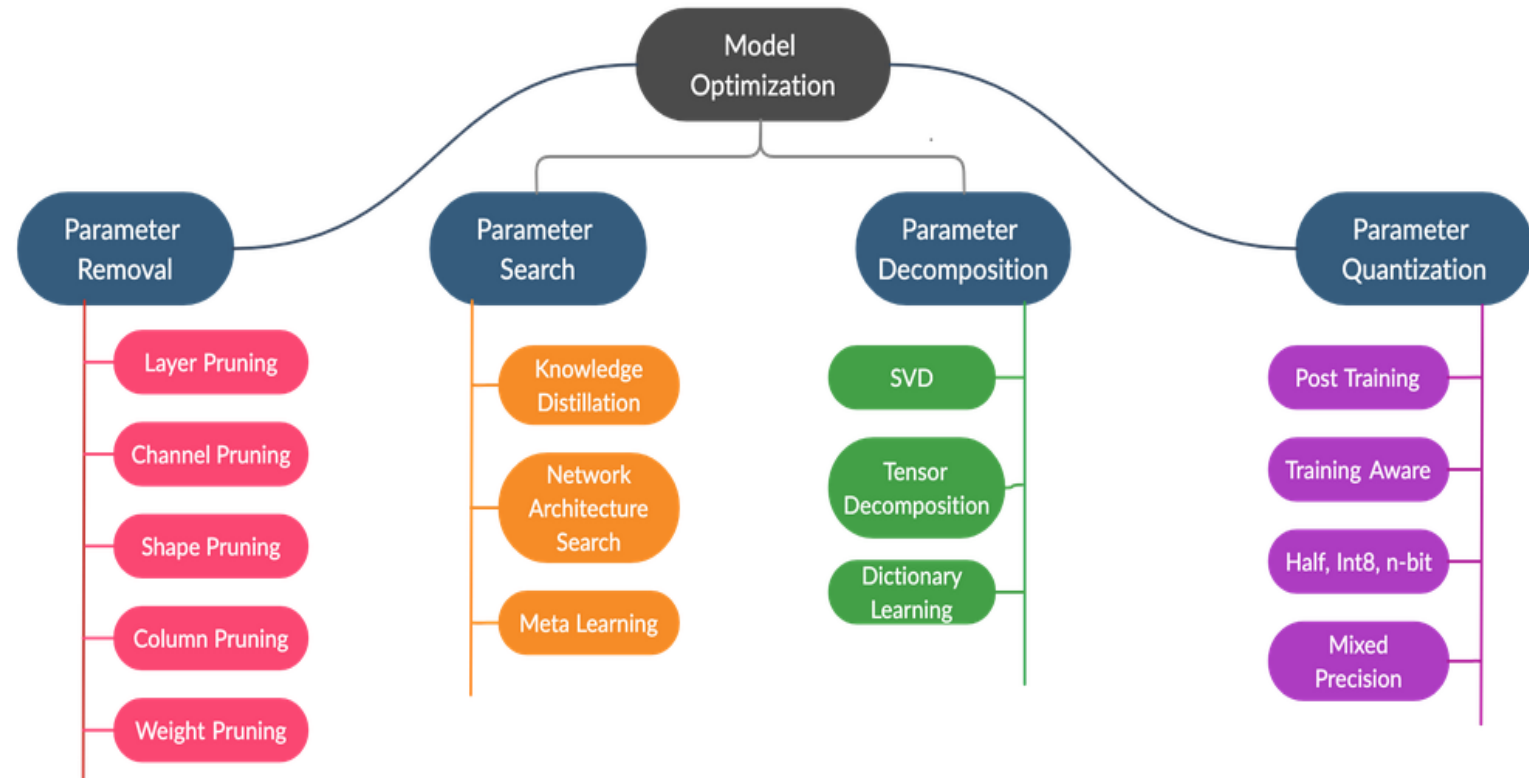
Abdullah, Muhammad  
Kakumanu, Ram Saran

### 01 Model Optimization: Quantization, Pruning and Approximate Computing

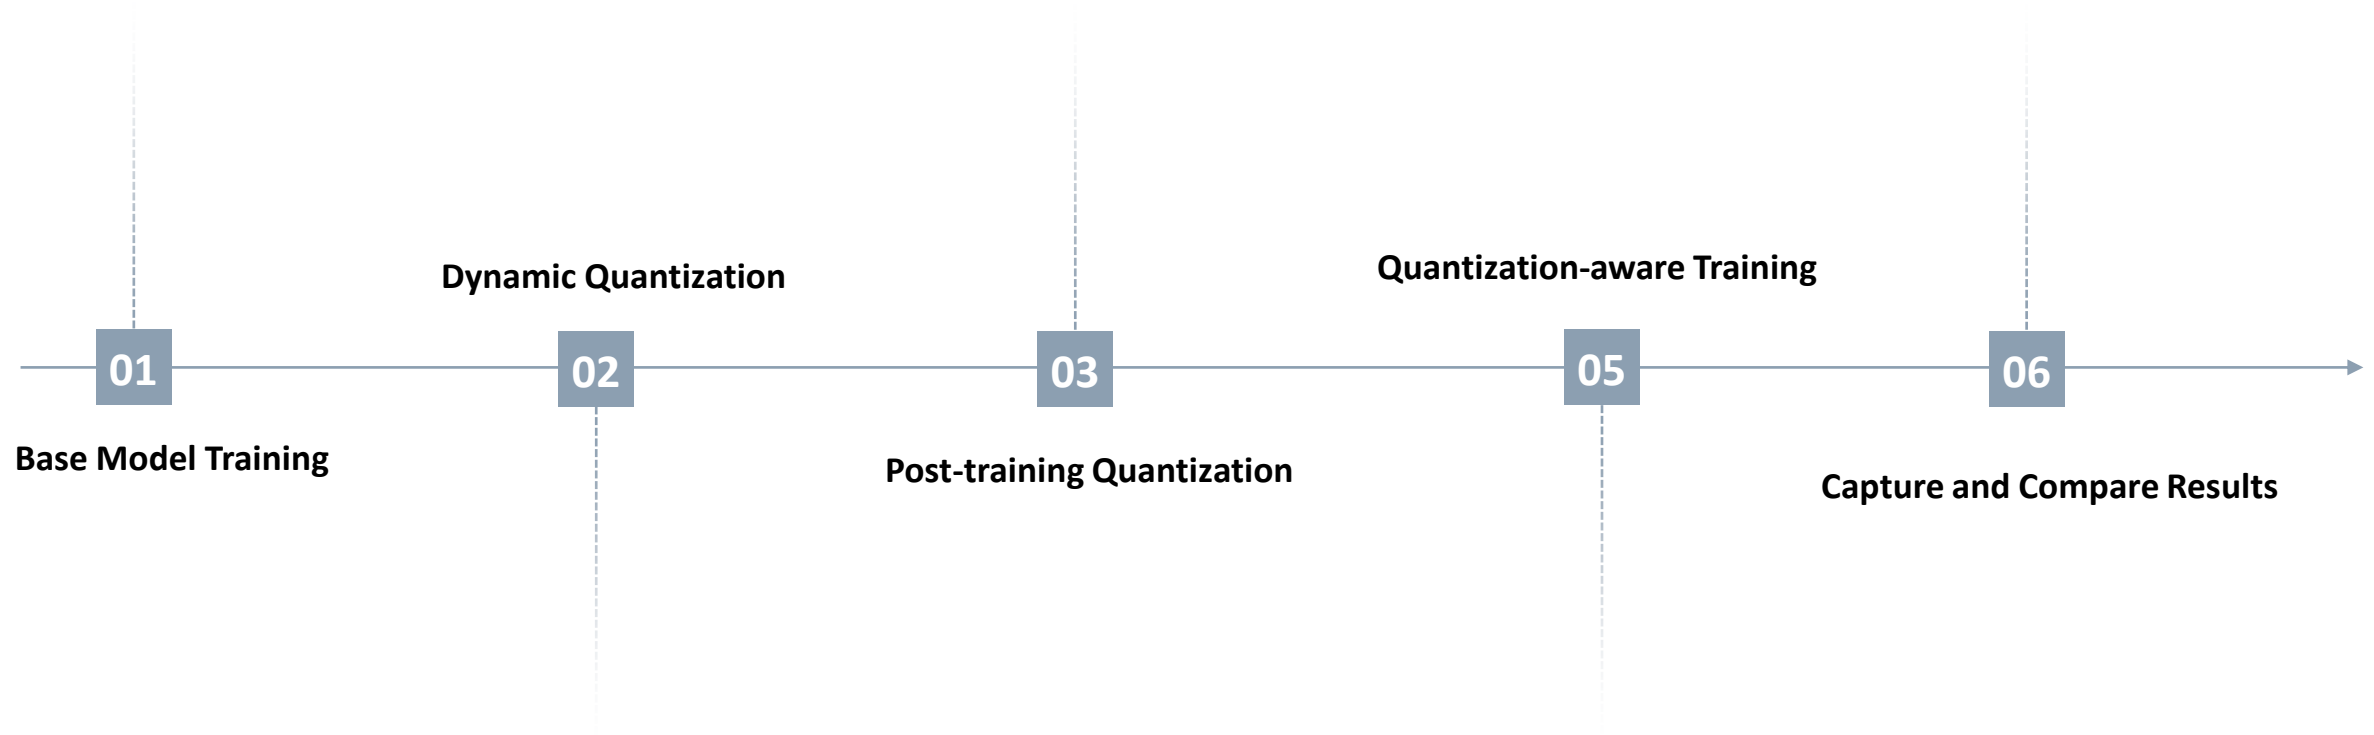
### 02 CIFAR datasets and ResNet Models

### 03 Quantization Experiments and Results

### 04 Pruning Experiments and Results




<https://www.edge-ai-vision.com/2020/09/dnn-model-optimization-series-part-i-whats-the-drill/>



Where bins for conversion of FP32 to Int8 are defined

### What are ResNet model?

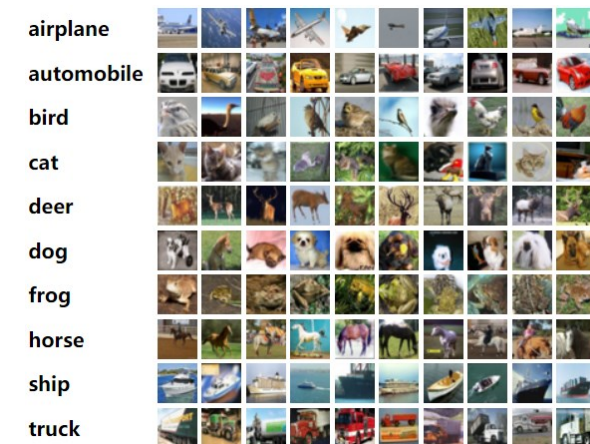
- DNN architecture widely used in computer vision tasks
- Uses concept of residual connections
- Addresses vanishing gradient problem



[https://www.researchgate.net/figure/Original-ResNet-18-Architecture\\_fig1\\_336642248](https://www.researchgate.net/figure/Original-ResNet-18-Architecture_fig1_336642248)

### What are CIFAR dataset?

- Benchmark DL models
- 60,000 32x32 color images
  - CIFAR10: 10 classes, with 6,000 images per class
  - CIFAR100: 100 classes, with 600 images per class



<https://www.cs.toronto.edu/~kriz/cifar.html>

### Dynamic Quantization

- Float32 x “Scalar Factor” = Rounded to nearest “Int8” → Dynamically at runtime
- Weights are known → Quantized before using
- Activations → Quantized on the fly (Before using in Activation Layers)
- Scaling factor adjusted based on input data
- Least performant quantization technique

```
import torch.quantization
quantized_model = torch.quantization.quantize_dynamic(model,
                                                         {torch.nn.Linear}, dtype=torch.qint8)
```

### Post Training Quantization

- Int8 Memory Access +
- Fine-tuning step between Model Completion and Inference
- Feed data batches → Distributions of different Activations → “Determines the bins”
- Right technique for medium-to-large-sized models

```
model_fp32 = CustomModelclass()
model_fp32.eval()

model_fp32.qconfig = torch.quantization.get_default_qconfig('fbgemm')

# fusing different layers into one
model_fp32_fused = torch.quantization.fuse_modules(
    model_fp32, [['conv', 'relu']])

# inserting observer module
model_fp32_prepared = torch.quantization.prepare(model_fp32_fused)

# quantization algorithm calibration using some data
model_fp32_prepared(input_fp32_model)

model_int8 = torch.quantization.convert(model_fp32_prepared)
output = model_int8(input_fp32)
```

### Quantization Aware Training

- Int8 mimic FP32 → “FakeQuantile”
- Training → FP32 only
- Best performance when compared to the other two methods
- Increased training time

```
model_fp32.train()
model_fp32.qconfig = torch.quantization.get_default_qat_qconfig('fbgemm')

model_fp32_fused = torch.quantization.fuse_modules(model_fp32,
                                                    [['conv', 'bn', 'relu']])
model_fp32_prepared = torch.quantization.prepare_qat(model_fp32_fused)

# calibration
training_loop(model_fp32_prepared)

model_fp32_prepared.eval()
model_int8 = torch.quantization.convert(model_fp32_prepared)
```

### Accuracy

- Measure of correctness
- Ratio of correct predictions
- Indicates model performance

### Inference Time

- Time taken to make predictions
- Measures runtime efficiency
- Need for real-time applications

### Training Time

- Duration required to train a model
- Determines practicality of model
- Parameters influence performance

### Model Size

- Amount of memory
- Occupied by the parameters
- Measured in bytes or megabytes

### Power Consumption

- Amount of electrical consumption during operation
- Determines energy efficiency of model

### Quantization Error

- Discrepancy from original model
- Arises due to the loss of precision
- Minimized to maintain accuracy



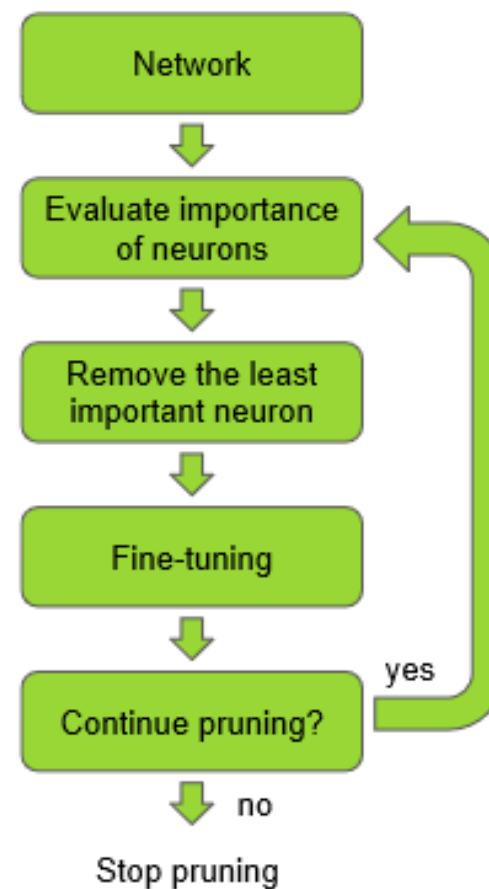
# DNN Quantization

Results: ResNet18 – CIFAR10 vs. CIFAR100

ResNet18 - Metrics	Base Model		Dynamic Quantization		Post Training Quantization		Quantization Aware Training	
	CIFAR10	CIFAR100	CIFAR10	CIFAR100	CIFAR10	CIFAR100	CIFAR10	CIFAR100
Accuracy	74.00%	42.79%	69.35%	35.36%	70.89%	37.63%	74.63%	42.79%
Model Size (in MB)	42.729	42.905	10.787	10.834	10.704	10.749	42.852	43.032
Training Time (in s)	1952.085	1866.672	1952.08	1866.672	1952.08	1866.672	3286.48	3307.339
Inference Time (in ms)	22.447	22.161	10.933	27.890	10.546	10.173	10.486	10.011
Quantization Error	NA	NA	0.790	0.976	0.791	0.975	0.253	0.572

- Trade-off between “model size and inference time” with accuracy
- As complexity of model/data increases – Quantization affects the Accuracy

- 01 Introduction
- 02 Terminology
- 03 Lottery Ticket Hypothesis
- 04 Pruning in Pytorch
- 05 Results
- 06 Appendix: A Note on HPC Parallel Experiments



<https://arxiv.org/pdf/1611.06440.pdf>

### Local Pruning

- Prune each layer/specified layer by a certain pruning ratio

### Global Pruning

- Prune weights globally. Some layers can have very high pruning ratio and some can have very low

### One-Shot Pruning

- Prune weights according pruning ratio all at once

### Iterative Pruning

- Prune progressively in iterations until reaching given pruning ration

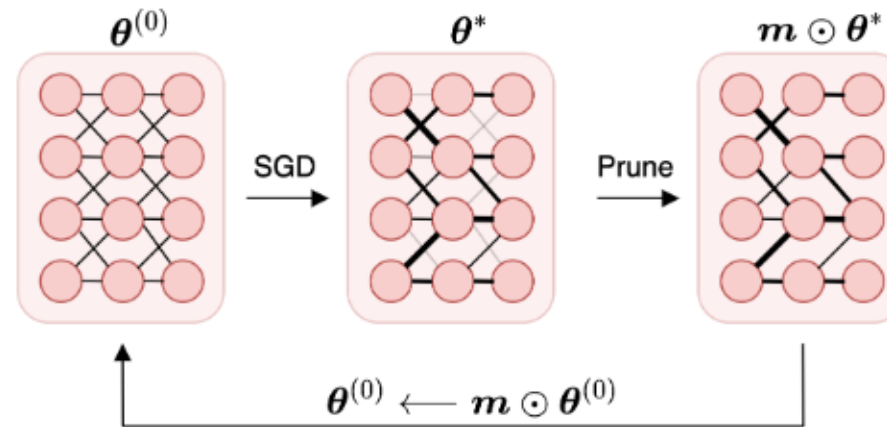
### Unstructured Pruning

- Remove connections without any pattern or constraints (Sparsity)

### Structured Pruning

- Remove entire structures or groups of connections while maintaining a certain pattern or structure.

- Sparse architectures after pruning are difficult to train from the start
- Dense, randomly initialized networks contain subnetworks which can achieve similar performance
- Copy weights of subnetworks from original weights (winning lottery ticket) after pruning
- Winning tickets learn faster than original network
- Prune FC layers (MLP) and conv layers (Vgg, Resnet)



Source: [https://www.researchgate.net/publication/368753812\\_Random\\_Teachers\\_are\\_Good\\_Teachers](https://www.researchgate.net/publication/368753812_Random_Teachers_are_Good_Teachers)

```
1  from torchvision.models import LeNet
2  from torch.nn.utils import prune
3
4  # Load Network
5  model = LeNet()
6
7  # Select layer you want to Prune
8  module = model.conv1
9
10 # Check current parameters for this layer
11 print(list(module.named_parameters()))
12
13 # Check buffer parameters for this layer
14 print(list(module.named_buffers()))
15
16 # Prune the layer by randomly making 30% weights zero
17 prune.random_unstructured(module, name="weight", amount=0.3)
18
19 # Pruning mask is stored in buffers names as 'weight_mask'
20 print(list(module.named_buffers()))
21
22 # A forward prehook is created
23 print(module._forward_pre_hooks)
24
25 # This is new pruned weights
26 print(module.weight)
27
```

[('weight', Parameter containing:  
tensor(...),  
device='cuda:0', requires\_grad=True)),  
  
('bias', Parameter containing:  
tensor(...),  
device='cuda:0', requires\_grad=True))]

```
1  from torchvision.models import LeNet
2  from torch.nn.utils import prune
3
4  # Load Network
5  model = LeNet()
6
7  # Select layer you want to Prune
8  module = model.conv1
9
10 # Check current parameters for this layer
11 print(list(module.named_parameters()))
12
13 # Check buffer parameters for this layer
14 print(list(module.named_buffers()))
15
16 # Prune the layer by randomly making 30% weights zero
17 prune.random_unstructured(module, name="weight", amount=0.3)
18
19 # Pruning mask is stored in buffers names as 'weight_mask'
20 print(list(module.named_buffers()))
21
22 # A forward prehook is created
23 print(module._forward_pre_hooks)
24
25 # This is new pruned weights
26 print(module.weight)
27
```

```
[('weight_mask',
 tensor([[1, 0, 1],...],
 device='cuda:0'))]
```

```
1  from torchvision.models import LeNet
2  from torch.nn.utils import prune
3
4  # Load Network
5  model = LeNet()
6
7  # Select layer you want to Prune
8  module = model.conv1
9
10 # Check current parameters for this layer
11 print(list(module.named_parameters()))
12
13 # Check buffer parameters for this layer
14 print(list(module.named_buffers()))
15
16 # Prune the layer by randomly making 30% weights zero
17 prune.random_unstructured(module, name="weight", amount=0.3)
18
19 # Pruning mask is stored in buffers names as 'weight_mask'
20 print(list(module.named_buffers()))
21
22 # A forward prehook is created
23 print(module._forward_pre_hooks)
24
25 # This is new pruned weights
26 print(module.weight)
27
```


```
[OrderedDict(
  [(0, <torch.nn.utils.prune.RandomUnstructured
    object at 0x7f0749753c70>)]]]
```

```
1  from torchvision.models import LeNet
2  from torch.nn.utils import prune
3
4  # Load Network
5  model = LeNet()
6
7  # Select layer you want to Prune
8  module = model.conv1
9
10 # Check current parameters for this layer
11 print(list(module.named_parameters()))
12
13 # Check buffer parameters for this layer
14 print(list(module.named_buffers()))
15
16 # Prune the layer by randomly making 30% weights zero
17 prune.random_unstructured(module, name="weight", amount=0.3)
18
19 # Pruning mask is stored in buffers names as 'weight_mask'
20 print(list(module.named_buffers()))
21
22 # A forward prehook is created
23 print(module._forward_pre_hooks)
24
25 # This is new pruned weights
26 print(module.weight)
27
```


```
tensor(...,
device='cuda:0', grad_fn=<MulBackward0>)
```




Parallel Coordinates View<sup>4</sup>



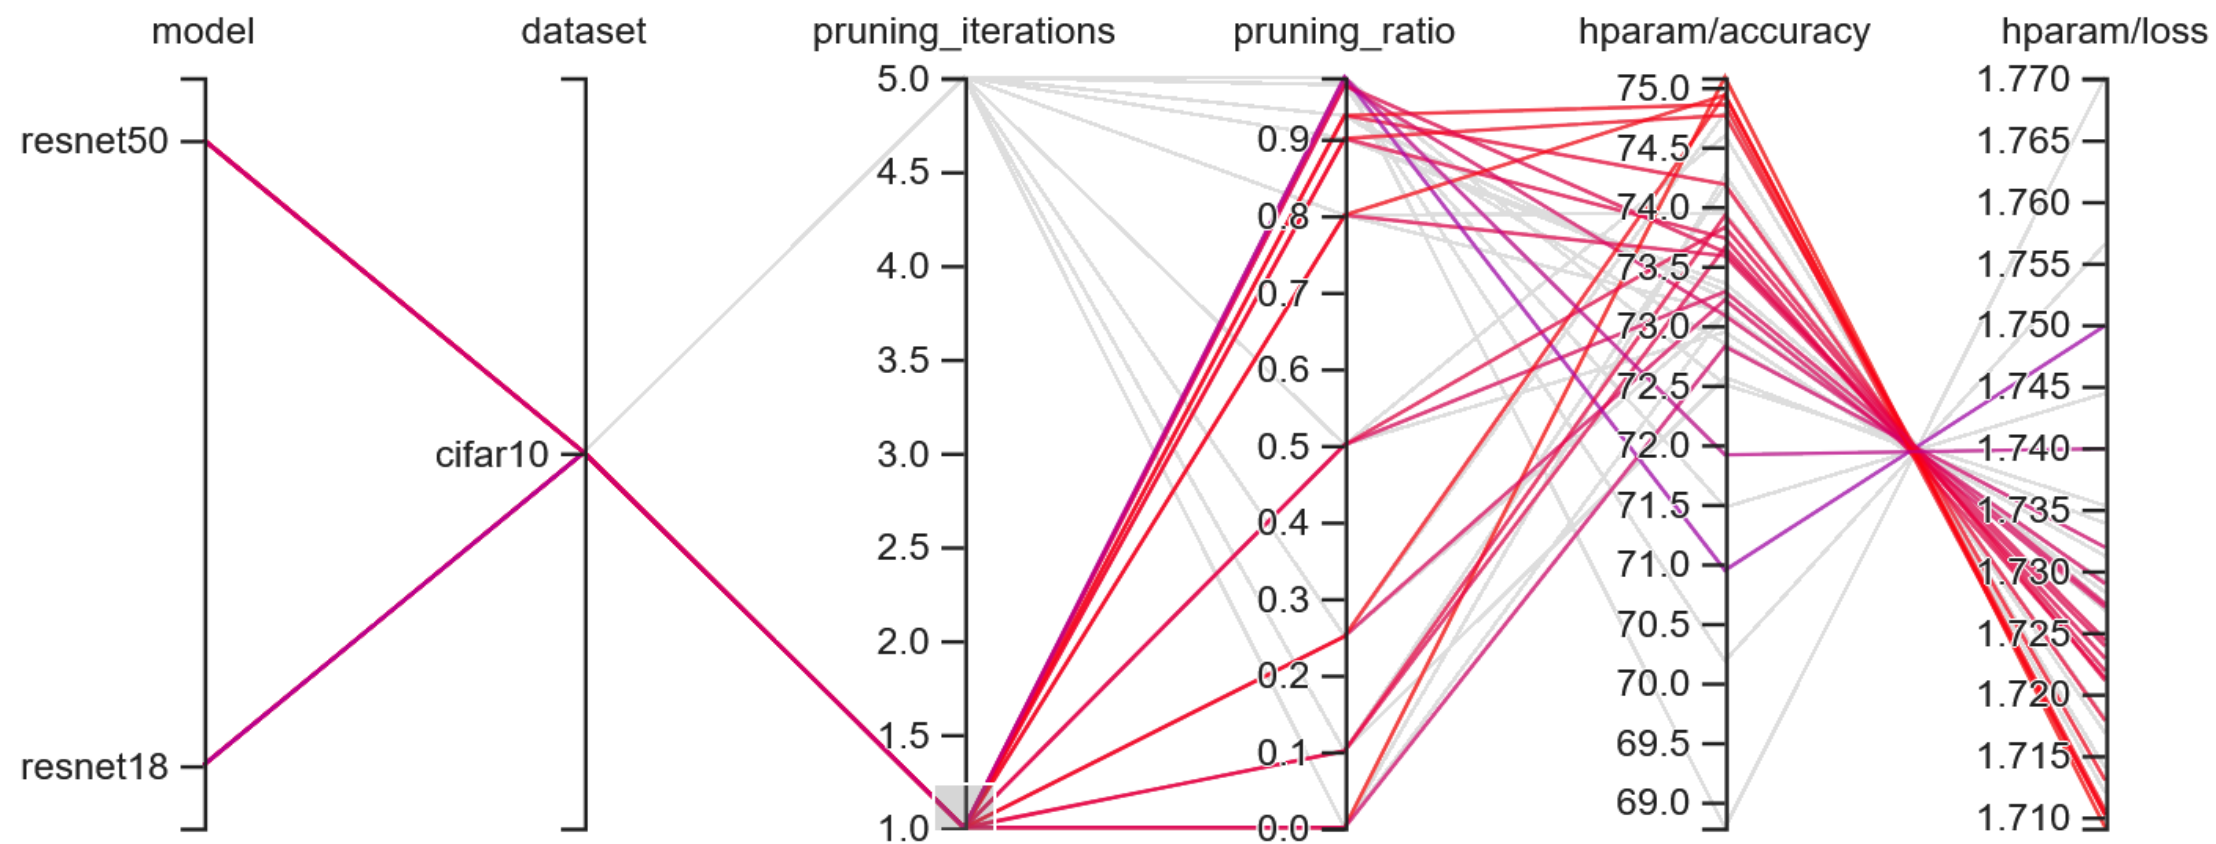
### Experiments with lowest accuracy




Experiments with highest accuracy



### Experiments with single shot pruning




### Experiments with iterative pruning



Cifar10 fine tuning perform better than random initialization of weights

Accuracy/validation





- Number of parameters in DNNs are exploding (175 billion in GPT3<sup>3</sup>)
- As use cases of DNNs are increasing, deploying them on edge devices with real-time performance poses a great challenge.
- Model optimization strategies like Quantization, Pruning, Approximate Computing etc are used to make DNNs deployable on resource-constrained devices.
- There is no one-size-fits-all solution. You have to perform experiments on your Dataset, Network to find best fit.

# Appendix: A Note on HPC Parallel Experiments

GNU *parallel* command<sup>5</sup>

```
$ parallel_exp.sh
1  #!/bin/bash
2  # Parallel Jobs
3  N_JOBS=10
4  ARGS="-P$N_JOBS --header :"
5
6  # Uncomment this line for dry run
7  #ARGS="--dry-run "$ARGS
8
9  # Experiment parameters
10 PROJECT='pruning_lottery_ticket_hypothesis'
11
12 MAX_EPOCHS=(50)
13 PRUNE_ITERS=(1 5)
14 PRUNE_METHODS=('l1')
15 PRUNE_RATIOS=(0 0.1 0.25 0.5 0.8 0.9 0.93 0.97 0.98)
16 REINITIALIZES=('false')
17 RANDOM_STATES=(1)
18 DATASETS=('cifar10')
19 MODELS=('resnet18' 'resnet50')
20
```

All parameters with different possible values are defined in a bash script first

```
20
21 parallel $ARGS \
22     sbatch \
23         --job-name=$PROJECT \
24         $(echo --export=dataset={dataset},\
25             model={model},\
26             epochs={max_epochs},\
27             pruning_iterations={prune_iter},\
28             pruning_method={prune_method},\
29             pruning_ratio={prune_ratio},\
30             seed={random_state},\
31             weight_reinit={reinitialize} | tr -d '[:space:]') \
32     run-job.sh \
33         ::: max_epochs "${MAX_EPOCHS[@]}" \
34         ::: prune_iter "${PRUNE_ITERS[@]}" \
35         ::: prune_method "${PRUNE_METHODS[@]}" \
36         ::: prune_ratio "${PRUNE_RATIOS[@]}" \
37         ::: random_state "${RANDOM_STATES[@]}" \
38         ::: reinitialize "${REINITIALIZES[@]}" \
39         ::: dataset "${DATASETS[@]}" \
40         ::: model "${MODELS[@]}" \
```

*parallel* command is used to submit jobs using SLURM<sup>6</sup> based *sbatch* command. ::: separator is used by parallel command to iterate over variables.



1. [https://pytorch.org/tutorials/intermediate/pruning\\_tutorial.html](https://pytorch.org/tutorials/intermediate/pruning_tutorial.html)
2. The Lottery Ticket Hypothesis: Finding Sparse, Trainable Neural Networks (<https://arxiv.org/abs/1803.03635>)
3. <https://en.wikipedia.org/wiki/GPT-3>
4. [https://en.wikipedia.org/wiki/Parallel\\_coordinates](https://en.wikipedia.org/wiki/Parallel_coordinates)
5. [https://www.gnu.org/software/parallel/parallel\\_tutorial.html](https://www.gnu.org/software/parallel/parallel_tutorial.html)
6. <https://slurm.schedmd.com/documentation.html>

## Notable Mentions for Programming help

- [https://github.com/facebookresearch/open\\_lth](https://github.com/facebookresearch/open_lth)
- [https://github.com/jankrepl/mildlyoverfitted/tree/master/github\\_adventures/lottery](https://github.com/jankrepl/mildlyoverfitted/tree/master/github_adventures/lottery)

---

# Thank You For Your Attention!